

# TD Learning

Shusen Wang

# Outline

1. Mathematically derive the TD target.
2. TD learning algorithms:
  - Sarsa algorithm for learning  $Q_{\pi}$ .
  - Q-learning algorithm for learning  $Q^*$ .
3. Multi-step TD target.


**Derive TD Target**

# Discounted Return

Definition of discounted return:

$$\begin{aligned} \bullet U_t &= R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \gamma^4 \cdot R_{t+4} + \dots \\ &\quad \underbrace{\hspace{10em}} \\ &= \gamma \cdot (R_{t+1} + \gamma \cdot R_{t+2} + \gamma^2 \cdot R_{t+3} + \gamma^3 \cdot R_{t+4} + \dots) \end{aligned}$$

# Discounted Return


- $$U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \gamma^4 \cdot R_{t+4} + \dots$$
$$= R_t + \gamma \cdot (R_{t+1} + \gamma \cdot R_{t+2} + \gamma^2 \cdot R_{t+3} + \gamma^3 \cdot R_{t+4} + \dots)$$


$= U_{t+1}$

# Discounted Return

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}$ .

- $$\begin{aligned} U_t &= R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot R_{t+3} + \gamma^4 \cdot R_{t+4} + \dots \\ &= R_t + \gamma \cdot (R_{t+1} + \gamma \cdot R_{t+2} + \gamma^2 \cdot R_{t+3} + \gamma^3 \cdot R_{t+4} + \dots) \end{aligned}$$


$$= U_{t+1}$$

# Derive TD Target

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}$ .

- Assume  $R_t$  depends on  $(s_t, A_t, s_{t+1})$ .
- $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t]$

# Derive TD Target

Identity:  $U_t = R_t + \gamma \cdot U_{t+1}$ .

- Assume  $R_t$  depends on  $(s_t, a_t, s_{t+1})$ .

- $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t]$   
=  $\mathbb{E}[R_t + \gamma \cdot U_{t+1} | s_t, a_t]$   
=  $\mathbb{E}[R_t | s_t, a_t] + \gamma \cdot \mathbb{E}[U_{t+1} | s_t, a_t]$



It is taken w.r.t.  $(s_{t+1}, a_{t+1}), (s_{t+2}, a_{t+2}), \dots$



# Derive TD Target

Identity:  $U_t = R_t + \gamma \cdot U_{t+1}$ .

- Assume  $R_t$  depends on  $(S_t, A_t, S_{t+1})$ .

- $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t]$   
 $= \mathbb{E}[R_t + \gamma \cdot U_{t+1} | s_t, a_t]$   
 $= \mathbb{E}[R_t | s_t, a_t] + \gamma \mathbb{E}[U_{t+1} | s_t, a_t]$

$$\mathbb{E}[U_{t+1} | s_t, a_t] = \mathbb{E}[Q_\pi(s_{t+1}, A_{t+1}) | s_t, a_t]$$

# Derive TD Target

Identity:  $U_t = R_t + \gamma \cdot U_{t+1}$ .

- Assume  $R_t$  depends on  $(S_t, A_t, S_{t+1})$ .

- $$\begin{aligned} Q_\pi(s_t, a_t) &= \mathbb{E}[U_t | s_t, a_t] \\ &= \mathbb{E}[R_t + \gamma \cdot U_{t+1} | s_t, a_t] \\ &= \mathbb{E}[R_t | s_t, a_t] + \gamma \mathbb{E}[U_{t+1} | s_t, a_t] \end{aligned}$$

$$\mathbb{E}[U_{t+1} | s_t, a_t] = \mathbb{E}[Q_\pi(s_{t+1}, A_{t+1}) | s_t, a_t]$$

$Q_\pi$  eliminates all the future states and actions from time  $t + 1$ .

# Derive TD Target

Identity:  $U_t = R_t + \gamma \cdot U_{t+1}$ .

- Assume  $R_t$  depends on  $(S_t, A_t, S_{t+1})$ .

- $$\begin{aligned} Q_\pi(s_t, a_t) &= \mathbb{E}[U_t | s_t, a_t] \\ &= \mathbb{E}[R_t + \gamma \cdot U_{t+1} | s_t, a_t] \\ &= \mathbb{E}[R_t | s_t, a_t] + \gamma \mathbb{E}[U_{t+1} | s_t, a_t] \end{aligned}$$

$$\mathbb{E}[U_{t+1} | s_t, a_t] = \mathbb{E}[Q_\pi(S_{t+1}, A_{t+1}) | s_t, a_t]$$



The expectation is taken w.r.t. only  $S_{t+1}$  and  $A_{t+1}$ .

# Derive TD Target

- Assume  $R_t$  depends on  $(S_t, A_t, S_{t+1})$ .
- $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t]$   
 $= \mathbb{E}[R_t + \gamma \cdot U_{t+1} | s_t, a_t]$   
 $= \mathbb{E}[R_t | s_t, a_t] + \gamma \cdot \mathbb{E}[U_{t+1} | s_t, a_t]$   
 $= \mathbb{E}[R_t | s_t, a_t] + \gamma \cdot \mathbb{E}[\underline{Q_\pi(S_{t+1}, A_{t+1})} | s_t, a_t].$

# Derive TD Target

**Identity:**  $Q_\pi(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_\pi(S_{t+1}, A_{t+1}) | s_t, a_t]$ , for all  $\pi$ .

- Assume  $R_t$  depends on  $(S_t, A_t, S_{t+1})$ .
- $Q_\pi(s_t, a_t) = \mathbb{E}[U_t | s_t, a_t]$   
     $= \mathbb{E}[R_t + \gamma \cdot U_{t+1} | s_t, a_t]$   
     $= \mathbb{E}[R_t | s_t, a_t] + \gamma \cdot \mathbb{E}[U_{t+1} | s_t, a_t]$   
     $= \mathbb{E}[R_t | s_t, a_t] + \gamma \cdot \mathbb{E}[Q_\pi(S_{t+1}, A_{t+1}) | s_t, a_t].$

# Derive TD Target

**Identity:**  $Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1}) | s_t, a_t]$ , for all  $\pi$ .

- We do not know the expectation.
- Approximate it using Monte Carlo (MC).

# Derive TD Target

**Identity:**  $Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1}) | s_t, a_t]$ , for all  $\pi$ .

$y_t$  is its MC approximation.

- Let  $(s_{t+1}, r_t)$  be an observation of  $(S_{t+1}, R_t)$ .
- Sample  $a_{t+1} \sim \pi(\cdot | s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot Q_{\pi}(s_{t+1}, a_{t+1})$ .

# Derive TD Target

**Identity:**  $Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1}) | s_t, a_t]$ , for all  $\pi$ .

$y_t$  is its MC approximation.

TD learning: Encourage  $Q_{\pi}(s_t, a_t)$  to approach  $y_t$ .



**Sarsa**

# Tabular Version

- We want to learn  $Q_{\pi}(s, a)$ .
- Suppose there are finite number of states and actions.
- Draw a table and learn the entries.

	Action $a_1$	Action $a_2$	Action $a_3$	Action $a_4$	...
State $s_1$					
State $s_2$					
State $s_3$					
⋮					

# Sarsa (tabular version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- Sample  $a_{t+1} \sim \pi(\cdot | s_{t+1})$ , where  $\pi$  is a policy function.
- TD target:  $y_t = r_t + \gamma \cdot Q_\pi(s_{t+1}, a_{t+1})$ .

	Action $a_1$	Action $a_2$	Action $a_3$	Action $a_4$	...
State $s_1$					
State $s_2$					
State $s_3$					
$\vdots$					

# Sarsa (tabular version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- Sample  $a_{t+1} \sim \pi(\cdot | s_{t+1})$ , where  $\pi$  is a policy function.
- TD target:  $y_t = r_t + \gamma \cdot Q_\pi(s_{t+1}, a_{t+1})$ .
- TD error:  $\delta_t = Q_\pi(s_t, a_t) - y_t$ .
- Update:  $Q_\pi(s_t, a_t) \leftarrow Q_\pi(s_t, a_t) - \alpha \cdot \delta_t$ .




make  $Q_\pi(s_t, a_t)$  closer to  $y_t$

# Value Network Version

- Approximate  $Q_{\pi}(s, a)$  by a value network,  $q(s, a|w)$ .



state  $s$



Value Network  
(parameterized by  $w$ )

$q(s, \text{"left"}; w)$

$q(s, \text{"right"}; w)$

$q(s, \text{"up"}; w)$

# Value Network Version

- Approximate  $Q_{\pi}(s, a)$  by a value network,  $q(s, a|\mathbf{w})$ .
- Note that the  $Q_{\pi}(s, a)$  and  $q(s, a|\mathbf{w})$  depends on  $\pi$ .
- $q$  is used as the critic who evaluates the actor. (Actor-Critic Method.)
- We seek to learn the parameter,  $\mathbf{w}$ .

# Sarsa (Value Network Version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- Sample  $a_{t+1} \sim \pi(\cdot | s_{t+1})$ , where  $\pi$  is a policy function.
- TD target:  $y_t = r_t + \gamma \cdot q(s_{t+1}, a_{t+1} | \mathbf{w})$ .
- TD error:  $\delta_t = q(s_t, a_t | \mathbf{w}) - y_t$ .
- SGD:  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \cdot \delta_t \cdot \frac{\partial q(s_t, a_t | \mathbf{w})}{\partial \mathbf{w}}$ .

# Sarsa: Recap

- **Goal:** Learn the action-value function  $Q_\pi$ .
- **Tabular version** (directly learn  $Q_\pi$ ).
  - There are finite state and actions.
  - Draw a table, and update the table using Sarsa.
- **Value network version** (function approximation).
  - Approximate  $Q_\pi$  by the value network  $q(s, a | \mathbf{w})$ .
  - Update the parameter,  $\mathbf{w}$ , using Sarsa.
  - Application: actor-critic method.



# Q-Learning

# Sarsa VS Q-Learning

- Sarsa is for training action-value function,  $Q_{\pi}(s, a)$ .
- TD target:  $y_t = r_t + \gamma \cdot Q_{\pi}(s_{t+1}, a_{t+1})$ .
- We used Sarsa for value network (critic).

# Sarsa VS Q-Learning

- Sarsa is for training action-value function,  $Q_{\pi}(s, a)$ .
- TD target:  $y_t = r_t + \gamma \cdot Q_{\pi}(s_{t+1}, a_{t+1})$ .
- We used Sarsa for value network (critic).
- Q-learning is for training the optimal action-value function,  $Q^*(s, a)$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .
- We used Q-learning for DQN.

# Derive TD Target

- We have proved that for all  $\pi$ ,

$$Q_{\pi}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi}(S_{t+1}, A_{t+1}) | s_t, a_t].$$

- If  $\pi$  is the optimal policy  $\pi^*$ , then

$$Q_{\pi^*}(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q_{\pi^*}(S_{t+1}, A_{t+1}) | s_t, a_t].$$

- We denote  $Q_{\pi^*}$  by  $Q^*$ .

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1}) | s_t, a_t].$

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1}) | s_t, a_t]$ .

- The action  $A_{t+1}$  is computed by

$$A_{t+1} = \underset{a}{\operatorname{argmax}} Q^*(S_{t+1}, a).$$

- Thus  $Q^*(S_{t+1}, A_{t+1}) = \max_a Q^*(S_{t+1}, a)$ .

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E}[R_t + \gamma \cdot Q^*(S_{t+1}, A_{t+1}) | s_t, a_t].$

- The action  $A_{t+1}$  is computed by

$$A_{t+1} = \underset{a}{\operatorname{argmax}} Q^*(S_{t+1}, a).$$

- Thus  $Q^*(S_{t+1}, A_{t+1}) = \max_a Q^*(S_{t+1}, a).$

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E} \left[ R_t + \gamma \cdot \max_a Q^*(S_{t+1}, a) \mid s_t, a_t \right].$

# Derive TD Target

**Identity:**  $Q^*(s_t, a_t) = \mathbb{E} \left[ R_t + \gamma \cdot \max_a Q^*(S_{t+1}, a) \mid s_t, a_t \right].$

# Derive TD Target

Identity:  $Q^*(s_t, a_t) = \mathbb{E} \left[ \underbrace{R_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)}_{\approx y_t} \mid s_t, a_t \right].$

- Let  $(s_{t+1}, r_t)$  be an observation of  $(S_{t+1}, R_t)$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a).$



# Q-Learning (tabular version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .

# Q-Learning (tabular version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \max_a Q^*(s_{t+1}, a)$ .

	Action $a_1$	Action $a_2$	Action $a_3$	Action $a_4$	...
State $s_1$					
State $s_2$					
State $s_3$					
$\vdots$					

# Q-Learning (tabular version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a)$ .
- TD error:  $\delta_t = Q^*(s_t, a_t) - y_t$ .
- Update:  $Q^*(s_t, a_t) \leftarrow Q^*(s_t, a_t) - \alpha \cdot \delta_t$ .


make  $Q^*(s_t, a_t)$  closer to  $y_t$

# DQN Version

- Approximate  $Q^*(s, a)$  by a value network,  $Q(s, a | \mathbf{w})$ .
- The network is called deep Q network (DQN).



state  $s$



DQN

(parameterized by  $\mathbf{w}$ )



$Q(s, \text{"left"}; \mathbf{w})$

$Q(s, \text{"right"}; \mathbf{w})$

$Q(s, \text{"up"}; \mathbf{w})$

# DQN Version

- Approximate  $Q^*(s, a)$  by a value network,  $Q(s, a|\mathbf{w})$ .
- The network is called deep Q network (DQN).
- It controls the agent by:  $a_t = \max_a Q(s_t, a|\mathbf{w})$
- We seek to learn the parameter,  $\mathbf{w}$ .

# Q-Learning (DQN Version)

- Observe  $(s_t, a_t, r_t, s_{t+1})$ .
- TD target:  $y_t = r_t + \gamma \cdot \max_a Q(s_{t+1}, a \mid \mathbf{w})$ .
- TD error:  $\delta_t = Q(s_t, a_t \mid \mathbf{w}) - y_t$ .
- Update:  $\mathbf{w} \leftarrow \mathbf{w} - \alpha \cdot \delta_t \cdot \frac{\partial q(s_t, a_t \mid \mathbf{w})}{\partial \mathbf{w}}$ .

# Q-Learning: Recap

- **Goal:** Learn the optimal action-value function  $Q^*$ .
- **Tabular version** (directly learn  $Q^*$ ).
  - There are finite state and actions.
  - Draw a table, and update the table using Q-learning.
- **DQN version** (function approximation).
  - Approximate  $Q^*$  by the value network  $Q(s, a | \mathbf{w})$ .
  - Update the parameter,  $\mathbf{w}$ , using Q-learning.

# Multi-Step Target



# Multi-Step Return

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}$

Replace  $U_{t+1}$  by  $R_{t+1} + \gamma \cdot U_{t+2}$

- It follows that  $U_{t+1} = R_{t+1} + \gamma \cdot U_{t+2}$ .

# Multi-Step Return

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}$

Replace  $U_{t+1}$  by  $R_{t+1} + \gamma \cdot U_{t+2}$

**Identity:**  $U_t = R_t + \gamma \cdot (R_{t+1} + \gamma \cdot U_{t+2})$

# Multi-Step Return

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}.$

Replace  $U_{t+1}$  by  $R_{t+1} + \gamma \cdot U_{t+2}$

**Identity:**  $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot U_{t+2}.$

# Multi-Step Return

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}.$

**Identity:**  $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot U_{t+2}.$

**Identity:**  $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot U_{t+3}.$

# Multi-Step Return

**Identity:**  $U_t = R_t + \gamma \cdot U_{t+1}.$

**Identity:**  $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot U_{t+2}.$

**Identity:**  $U_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+2} + \gamma^3 \cdot U_{t+3}.$

**Identity:**  $U_t = \sum_{i=0}^{m-1} \gamma^i \cdot R_{t+i} + \gamma^m \cdot U_{t+m}.$

# Multi-Step TD Targets

**Identity:**  $U_t = \sum_{i=0}^{m-1} \gamma^i \boxed{R_{t+i}} + \gamma^m \boxed{U_{t+m}}.$

- $m$ -step TD target for learning  $Q_\pi$  (i.e., Sarsa):

$$y_t = \sum_{i=0}^{m-1} \gamma^i \cdot \boxed{r_{t+i}} + \gamma^m \cdot \boxed{Q_\pi(s_{t+m}, a_{t+m})}.$$

# Multi-Step TD Targets

**Identity:**  $U_t = \sum_{i=0}^{m-1} \gamma^i \boxed{R_{t+i}} + \gamma^m \boxed{U_{t+m}}$

- $m$ -step TD target for learning  $Q_\pi$  (i.e., Sarsa):

$$y_t = \sum_{i=0}^{m-1} \gamma^i \cdot \boxed{r_{t+i}} + \gamma^m \cdot Q_\pi(s_{t+m}, a_{t+m}).$$

- $m$ -step TD target for learning  $Q^*$  (i.e., Q-learning):

$$y_t = \sum_{i=0}^{m-1} \gamma^i \cdot \boxed{r_{t+i}} + \gamma^m \cdot \boxed{\max_a Q^*(s_{t+m}, a)}$$

# Comparison

- One-step TD target for learning  $Q^*$  (or DQN):

$$y_t = r_t + \gamma \cdot \max_a Q^*(s_{t+1}, a).$$

- $m$ -step TD target for learning  $Q^*$  (or DQN):

$$y_t = \sum_{i=0}^{m-1} \gamma^i \cdot r_{t+i} + \gamma^m \cdot \max_a Q^*(s_{t+m}, a).$$

- If  $m$  is suitably tuned,  $m$ -step target works better than one-step target [1].

## Reference:

1. Hossel et al. [Rainbow: combining improvements in deep reinforcement learning](#). In AAAI, 2018.



# Summary

# Summary

1. Mathematically derived the TD target.
2. TD learning algorithms:
  - Sarsa algorithm for learning  $Q_\pi$  and value network (critic).
  - Q-learning algorithm for learning  $Q^*$  and DQN.
3. Multi-step TD target.

# Improvements for TD Learning

1. Experience replay (ER) and its variants (e.g., prioritized ER.)
  - Reuse experience.
  - Eliminate correlation.
2. Target network and double DQN.
  - Address the overestimation issue of Q-learning.
3. Multi-step TD target.

**Thank you!**